

From Courts to Comprehension: Can LLMs Make Judgments More Accessible?

Giovanni Pinna[†], Davide Tugnoli^{*}, Mara Bartole^{*}, Luca Manzoni[†], Andrea De Lorenzo^{*}

^{*}Department of Engineering and Architecture, University of Trieste, Trieste, Italy

[†]Department of Mathematics, Computer Science and Geosciences, University of Trieste, Trieste, Italy

Abstract—Legal regulations play a pivotal role in shaping citizens’ daily lives, yet their complexity often renders them inaccessible to those without specialized legal expertise. Recent advances in natural language processing (NLP) have shown promise in creating summaries of legal texts to enhance their comprehensibility. However, the effectiveness of these summaries, particularly when generated by Large Language Models (LLMs), has not been extensively evaluated among the general public audience—i.e. non-experts.

This study evaluates the capability of LLMs, specifically small open-source models and GPT-4o, in summarizing Italian legal judgments to make them more understandable for individuals without legal training. To assess the quality and comprehensibility of these summaries, participants were presented with a questionnaire containing comprehension questions formulated by legal experts. While this response doesn’t directly measure the summary’s quality, it serves as a strong indicator of the summary’s practical usefulness. The findings reveal that although these models are not yet fully capable of the task, they have demonstrated significant promise. However, the study also showed that while human-made summaries resulted in better comprehension and more accurate responses, they come at a higher cost compared to AI-generated summaries.

Index Terms—Legal Summarization, Natural Language Processing, Large Language Models, Human Evaluation, Legal Accessibility

I. INTRODUCTION

In a rule-of-law society, legal regulations profoundly impact citizens’ daily lives, making them a crucial element for individuals. However, understanding these laws poses significant challenges for those without specialized legal knowledge. This difficulty arises from the complex language and structure of legal texts, which are often lengthy and intricate [36]. Over the past year, LLMs have demonstrated remarkable performance, and these models are increasingly being utilized across various domains and tasks. To achieve the goal of summarizing and making legal texts more understandable, much research has used the latest NLP techniques, including LLMs. In the Italian legal environment, summaries known as *massime* condense the key reasons and decisions of judgments, allowing for quicker and more direct access to the main points. However, the *massime* are not intended for a general audience, but they are aimed at highlighting a particular aspect of a legal text. Despite their undeniable utility, their creation is labor-intensive and costly due to the extensive and complex nature of legal texts. This makes the prospect of employing LLMs to generate comprehensible summaries for non-experts highly appealing. Despite the promising results of previous studies [11] [21] [1],

the evaluation of the quality of summaries on the target audience of people without legal knowledge has not yet been explored. Evaluation of achievements is a crucial step, and automatic evaluation metrics are not always sufficient [8]. In fact, the use of LLMs to summarize legal texts raises important ethical questions. The ability of these systems to generate automated summaries of complex legal documents can increase the efficiency of legal work but requires careful consideration of their reliability and accountability.

In this study, we aim to evaluate the capabilities of LLMs in summarizing Italian legal texts to make them more easily understandable by the general public. In particular, we applied an *extrinsic* measurement, meaning we measured how reading a summary created by an LLM impacts text comprehension abilities compared to reading the original text or a summary made by a person. To do this, we administered a multiple-choice questionnaire, prepared by a legal expert, to people who could respond while having the text available. In this way the quality of the summaries is assessed indirectly by evaluating how well the summary helps end users understand what is written in the original text.

To support the results of our analysis, we conducted statistical tests, which showed the potential of LLMs in the legal field. Our results suggest that, while GPT-4o-generated summaries offer a faster and more cost-effective alternative to traditional *massime*, they fall short in accuracy and clarity. The small open-source models also lag behind, indicating a need for further development. Our analysis highlights the potential of LLMs in the legal domain but also underscores the necessity for improvements. Future research should focus on enhancing the context-handling capabilities of open-source models and integrating legal knowledge to improve summary accuracy. Additionally, involving legal experts in the evaluation process could provide deeper insights into the challenges of using LLMs for legal text summarization, ultimately making legal information more accessible to the general public.

II. RELATED WORK

In the field of automatic text summarization, there are two main methods: extractive and abstractive. Extractive methods focus on selecting and combining sentences from the original text, often resulting in fluent outputs since they use human-written content [17] [31]. Some extractive approaches use parsed sentence structures to remove noise and highlight

relevant content [22]. To advance towards abstractive summarization, researchers have introduced sentence compression techniques, which aim to reduce noise and condense information [15] [33] [3]. However, maintaining consistency in these summaries has been a difficult challenge. Some studies have focused on text-to-text generation methods using neural networks [26] [35] [18], for generating summaries from input documents [27] [32] [19] [34]. These methods train neural networks to automatically extract syntactic and semantic features from input text to generate summaries. However, they often provide limited control over summary attributes and are less tractable. Addressing these challenges is crucial for improving the effectiveness and versatility of text-to-text generation techniques. The recent interest in NLP methods has been fueled by the introduction of LLMs [29] [25]. These models, with their extensive parameters, exhibit learning capabilities and are specifically trained for various NLP tasks, including text-to-text generation and machine translation [23]. However, memory-efficient LLMs are limited by their low context length [9] and may still experience hallucinations [4].

The legal field governs the lives of every individual every day. Given its importance, there are many studies on the application of artificial intelligence techniques to this field dating back to the early 1970s. Among the pioneering works is Buchanan and Headrick’s article [7], which explores how AI can support legal functions. The authors imagined computers not only as repositories for legal data but also as tools for facilitating effective communication with lawyers, assisting in their daily tasks, and aiding in legal research. This perspective is reflected in numerous subsequent studies first of all the analysis in [2] illustrates how artificial intelligence (AI), particularly through NLP and machine learning (ML), is transforming the legal profession. With these insights, some recent research has aimed at even more ambitious goals, such as making legal texts more comprehensible to non-specialists. For example, Rahul C. Kore et al. [14] developed a solution that simplifies the understanding of extensive legal documents using NLP and ML techniques. They employ an extractive synthesis method, selecting key phrases from the original text while maintaining their integrity.

Recent advancements in LLMs offer additional prospects for the legal industry. In this sense, a possible approach is explored by Achena et al. in [1], who investigated summarization models to generate *massime* from rulings of the Italian Supreme Court and Constitutional Court. They compared different summarization models based on Bidirectional Encoder Representations from Transformers (BERT) [12], adapted to the legal domain, and evaluated them using the Recall-Oriented Understudy for Gisting Evaluation (ROUGE) metric [16]. The results highlight the difficulties in creating *massime*, compounded by differences in court judgments. Another important work was done by Martin et al. [21] examined the effectiveness of these models for tasks such as contract review, comparing them with traditional, human-based methods regarding accuracy, speed, and cost-effectiveness. This analysis, conducted within the legal systems of the United States and New Zealand,

which follow common law, revealed that LLMs outperform humans in speed and cost while maintaining comparable accuracy. These findings suggest a potential paradigm shift in the use of LLMs in future legal practice. Furthermore, Aniket Deroy et al. [11] investigated the applicability of automatic abstract synthesis models and LLMs for summarizing judgments of the Indian Supreme Court. They applied several legal domain-specific abstract synthesis models, such as Legal-Pegasus¹ and Legal-LED², alongside LLMs like ChatGPT and Davinci. The syntheses were evaluated using metrics such as ROUGE, Metric for Evaluation of Translation with Explicit ORDERing (METEOR) [5], and Bilingual Evaluation Understudy (BLEU) [24]. Their results indicated that while abstract models outperformed extractive models, they often contained inconsistent or hallucinated information [11].

III. OUR APPROACH

The purpose of this study is to evaluate the abilities of LLMs in synthesizing and summarizing Italian legal texts in a way that makes them more understandable to the general public. To achieve this goal, we prepared a questionnaire to evaluate users’ comprehension of four legal texts. The user is provided with a text without knowing its category (original text, *massima*, or generated summary) and can consult it while answering the questions. Subsequently, we can indirectly evaluate the quality of the summary by determining whether its reading led the user to identify the correct answer. More in detail, the texts submitted to each user are: (i) one full judgment, (ii) one *massima*, (iii) and two generated summaries. For each text, the user had to answer three multiple-choice questions with five possible answers, one of them being I don’t know. The questions were developed by a legal expert with access only to the original judgment. To avoid introducing bias, the expert was not exposed to either the *massime* or the generated summaries. The generated summaries and the *massime* have the same questions as the judgment from which they are derived.

It is important to notice that some questions could not be answered only based on the generated summary by LLMs or *massime*. By their nature, model summaries may have overlooked critical information present in the original text. *Massime*, instead, sometimes include only a few points from the original judgment: a *massima* encapsulates the principle of law established by a judgment, offering a concise view of the legal principles applied to a specific case. This makes *massime* useful for extrapolating key points of often complex and detailed judicial decisions, aiding understanding for legal professionals. As such, they might not contain all the information present in the original judgment and are tailored more specifically to the legal experts and not to the general public.

The creation of *massime* requires in-depth analysis of judgments to extract relevant legal principles, a task typically performed by specialized offices in the Italian legal system at

¹<https://huggingface.co/ksi319/legal-pegasus>

²<https://tinyurl.com/Legal-LED>

the great expense of resources. As already stated, *massime* are aimed at a different audience than the general public, we decided, however, to use them as a benchmark for the generated texts because they are the closest to a well-done summary and are easily available on a large scale.

IV. EXPERIMENTAL EVALUATION

We selected Gemma 2B and 7B [28], LLaMA 2 7B [30], Llamantino 7B [6], and GPT-4o³ for this task. We selected open-source models with a small number of parameters that can run on consumer hardware. To compare with the state-of-the-art, we also selected GPT-4o, since it is available at low cost to consumers. Since open-source models are not specifically trained for Italian text, and particularly not for Italian legal text, we performed a fine-tuning, using *massime*. To limit the length of the questionnaire, we selected only one of the open-source models after a first qualitative evaluation.

The experimental phase is structured as follows: (i) identification of judgments and their respective *massime*, (ii) fine-tuning of LLMs, (iii) qualitative evaluation of summaries generated by fine-tuned and non-fine-tuned LLMs, (iv) generation of legal judgment summaries, (v) assessment of users' understanding of the summaries through a questionnaire made available via a web application, (vi) collection and analysis of users' answers.

A. Data collection

In this research, we used judgments and *massime* issued by the Italian Constitutional Court⁴. This choice is mainly motivated by the availability of a public dataset⁵. The Italian Constitutional Court's powers include judging the disputes of constitutional legitimacy of laws and the conflict of attribution of powers to the State or between the State and the Regions.

The open data from the Constitutional Court includes a collection of pronouncements and their related *massime* from 1956 to the present day. In total, we collected 10 000 pairs of rulings and *massime* and used 95% of them for fine-tuning. From the remaining 5%, we extracted 25 pairs, ensuring an even coverage of the years from 2000 to 2023. This decision was prompted by the concern that using older judgments could involve outdated legal terminology, possibly confusing the reader.

B. LLMs

All the open-source models chosen for this study are available on HuggingFace⁶. For the fine-tuning of the LLMs, we used one booster node of the Leonardo supercomputer⁷ with four NVIDIA A100 GPUs with 64GB of VRAM each. We fine-tuned each model for 3 epochs, with the Paged AdamW optimizer [20] with a learning rate of 0.0002, and a weight decay of 0.03. To optimize the model's ability to understand

and respond to the instructions, we employed the official tags provided by the model developers, as listed in Table I. For generating summaries, we employed two types of prompts: the system prompt and the specific instruction for generating the summary. The system prompt for both LLaMA 2 and LLaMAntino is the Italian translation of the standard one provided alongside the model. We observed better results with the Italian prompt because the model was more inclined to follow our instructions and provide responses in Italian. The instructions for generating the summary are the following:

Rispondi in italiano e fai il riassunto del seguente testo.

For the Gemma model, there is no system prompt. However, the model does follow a direct prompt for the task of summarizing judgments:

Sei un modello esperto di leggi, sentenze e ordinanze italiane. Di seguito è riportata una sentenza o una legge italiana, il tuo compito è quello di riassumerla. Scrivi una risposta che soddisfi adeguatamente questa richiesta.

The full text of all prompts in both English and Italian is available in the project repository <https://github.com/giovannipinna96/LawGPT/tree/master/prompt>.

TABLE I
TAGS USED FOR FINE-TUNING AND INFERENCE OF LLMs.

Model	Tag used
Gemma	<start_of_turn>user\n request <end_of_turn>\n<start_of_turn>model\n
LLaMA 2	<s>[INST] <<SYS>>\n System Prompt \n</SYS>>\n\n Before the sentence
LLaMAntino	<s>[INST] <<SYS>>\n System Prompt \n</SYS>>\n\n Before the sentence

It is important to note that longer judgments may exceed the context window limit of certain models. For LLaMA 2, which has a context window of 4k tokens, 21 of the 25 judgments exceed the token limit. For Gemma, which has a longer context window of 8k tokens, this happens only 10 out of 25 times. In these circumstances, sentences are truncated to the context window length. In contrast, GPT-4o has a context window of about 128k tokens and can handle all judgments without truncation. The qualitative analysis of the results produced by the open-source models (both with and without fine-tuning) gave us the following results: (i) Gemma 2B and 7B showed poor performance in Italian. The generated responses often included HTML tags or failed to address queries correctly, both in their standard and fine-tuned versions; (ii) LLaMA 2 responded consistently in English despite using Italian prompts; (iii) the fine-tuned LLaMA 2 had better performance in generating legal judgment summaries in Italian, making it the preferred candidate for comparison with GPT-4o; (iv) LLaMAntino while demonstrating similar performance to the fine-tuned LLaMA 2, occasionally included tags related to the system prompt in the responses. Due to these observations, the fine-tuned LLaMA 2 (FT-LLaMA 2 from now on) was selected

³<https://openai.com/index/hello-gpt-4o/>

⁴<https://www.cortecostituzionale.it/>

⁵https://dati.cortecostituzionale.it/Scarica_i_dati/Scarica_i_dati/

⁶<https://huggingface.co>

⁷<https://leonardo-supercomputer.cineca.eu/>

as the best overall performing model among the open-source ones.

C. Collection of questionnaire answers

To collect the answers to the questionnaires, we developed a web application⁸ accessible to the study participants. Upon accessing the web page, participants are greeted with a brief introductory guide and an initial survey, which collects preliminary information about the participant, including educational background and legal knowledge. After the initial survey, the participant is subjected to four legal texts of different categories presented in random order: an original judgment, a *massima*, a summary generated by an open source model, and a summary generated by GPT-4o. Each text refers to a different judgment to ensure that all questions are independent from each other. The user is also allowed to skip to the next text without completing all the answers for the current text.

At the end of the survey period, a total of 75 individuals, all of whom had voluntarily chosen to participate, were included in the study. Of these 75 volunteers, 25% said they had some legal knowledge while the remaining 75% said they had no legal knowledge.

V. RESULTS

The analysis of the obtained answer will focus on four different aspects: (i) percentage of correct answers, presented in Figures 1 and 2; (ii) percentage of incorrect answers (excluding I don't know), presented in Figs 3 and 4; (iii) percentage of I don't know answers, shown in Figures 5 and 6; (iv) percentage of skipped answers, presented in Figure 5; (v) reading and response times, shown in Figure 9.

In all cases, we show the results considering all participants and a subdivision between participants with and without legal knowledge. Statistical tests were also performed to see if a difference was present between the different text categories.

With the contingency table resulting from the data of the first four cases we use the Chi-squared test [10] to decide if the responses of people with and without legal knowledge were different (p-value of $4 \cdot 10^{-5}$). We then compared the influence of the different text types in the two sub-populations of people with and without legal knowledge using the Chi-squared test with Bonferroni-Holm correction [13].

With a significance level of $\alpha = 0.05$, the following differences were measured: 1) for participants with legal knowledge, there is a significant difference in the responses between the following pairs of text type: GPT4-o summaries and original judgments ($p = 0.011$), *massime* and original judgments ($p = 0.011$), and *massime* and FT-LLaMA2 summaries ($p = 0.011$). The corrected p-value for the comparison between GPT4-o and FT-LLaMA2 summaries is just above the significance threshold at 0.054; 2) for participants without legal knowledge, there is a significant difference in the responses between the same pairs of text type: GPT4-o summaries and original judgments ($p = 0.038$), *massime* and

original judgments ($p = 0.001$), and *massime* and FT-LLaMA 2 summaries ($p = 0.038$).

While the Chi-squared test on the contingency table can detect the differences in the responses, in the following we will analyse each of the four categories (correct, incorrect, I don't know, and skipped answers) in detail.

Percentage of correct answers

The results presented in Figure 1 show some general trends. First of all, in almost all cases people with legal knowledge answer correctly more frequently than people without it. The summary by FT-LLaMA2 is, for all categories, the one with the smallest percentage of correct answers, followed by the original judgment. Notice how, in this case, there is only a 3% difference between people with and without legal knowledge. Better than the original judgments are the results obtained with the summary provided by GPT-4o, with an increase of 4% for people with legal knowledge and 6% for people without. The best results are obtained with the *massime*, where the aggregate results are of 45% of correct answers, increasing to 58% for participants with legal knowledge, the only case where the percentage correct answers climb over 50%. The *massime* also represents the case where the effect given by legal knowledge is more pronounced, with a difference of 17%.

A more detailed comparison is possible by observing the results presented in Figure 2. There, the effect of preexisting legal knowledge on the distribution of correct answers is visible. For FT-LLaMA2 and the original judgment more than 40% of the participants answered correctly 0 out of 3 questions. For GPT-4o almost 50% of the participants answered correctly 1 out of 3 questions, while for *massime* almost 40% of the people answered correctly all questions. The pattern is very different in the case of people without any legal knowledge. There, around 40% of all people were unable to answer correctly even one question for the LLM-generated summaries and the original judgment. The best results are still given by the *massime* but in this case the distribution peaks at 1 out of 3 answers, with less than 15% of the people being able to answer correctly all three questions.




Fig. 1. Percentage of correct answers for each type of text for all participants (all), for participants with legal knowledge (with), and for people without legal knowledge (w/o)

⁸<https://web.archive.org/web/20240813090453/https://sentence.inginf.units.it/>




Fig. 2. Distribution of the number of correct answers according to the type of text for all participants (left), participants with preexisting legal knowledge (centre), and participants without legal knowledge (right).

Percentage of incorrect answers

Recall that there were three possible kinds of answers, one of them being *I don't know*, so it is interesting to explore if a low percentage of correct answers is accompanied by incorrect answers or by the perceived absence of the necessary information for answering the question. As it is possible to observe in Figure 3, the lowest percentage of incorrect answers is with *massime* and the original judgments, independently from the legal expertise of the participants. A higher percentage of incorrect answers is present for the summaries produced by FT-LLaMA2, but with similar percentages for both people with and without preexisting legal knowledge. There is, however, a large difference in the case of the summaries generated by GPT-4o, where more than 50% of the people with legal knowledge provided incorrect answers, compared to only 35% for the other participants.

Similar results are also visible when exploring the distribution of the number of incorrect answers in Figure 4. For participants without legal knowledge, the four distributions (depending on the type of text) are all similar. In contrast, for participants with legal knowledge, the distributions are quite distinct. It is possible to observe that the *massime* provides a minimal amount of incorrect answers (38.9%). At the same time, even the original judgments have less than 20% of zero incorrect answers. GPT-4o, in contrast, has a distribution peaking (with 61% of the participants) at 2 incorrect answers out of 3.

Percentage of *I don't know* and skipped answers

An analysis of the *I don't know* and skipped answers is necessary to complement the results on correct and incorrect answers. The percentage of *I don't know* answers is presented in Figure 5 and their distribution in Figure 6. The percentage of skipped answers (i.e., all questions related to a particular text were skipped without giving a response) is presented in Figure 7. We can notice two main trends. The first is that people with legal knowledge are less prone to answering *I don't know* than participants without such knowledge independently from the type of text. The second trend is




Fig. 3. Percentage of incorrect answers for each type of text for all participants (all), for participants with legal knowledge (with), and for people without legal knowledge (w/o).

about the difference between LLM-generated summaries and *massime*. While the former have percentages to the original judgments, the latter have higher percentages of *I don't know* answers, possibly highlighting the fact that *massime* captures only some aspects of the original text.

The distribution of the answers confirms both trends. Again, participants with legal knowledge were less prone to answering *I don't know* compared with people without legal knowledge (for three of the four test categories none of the former answered *I don't know* in all three questions). Also, the distribution of answers for the *massime* is more skewed towards higher numbers for both categories of participants.

Finally, the percentages of skipped answers presented in Figure 7 show that the original judgments and the summaries generated via FT-LLaMA2 are the ones with a high percentage of skipped answers compared to *massime* and summaries generated via GPT-4o. As we will see in the next section, this is strongly correlated with the length of the texts.

Reading and response times

As stated above, the results can be influenced by the different lengths and time needed to read and understand the texts. The distribution of the lengths of the different categories of




Fig. 4. Distribution of the number of incorrect answers according to the type of text for all participants (left), participants with preexisting legal knowledge (centre), and participants without legal knowledge (right).




Fig. 5. Percentage of I don't know answers for each type of text for all participants (all), for participants with legal knowledge (with), and for people without legal knowledge (w/o)

texts is presented in Figure 8. As it is possible to observe, the shorter texts are the *massime* and the summaries generated by GPT-4o, all below 10^3 words. The summaries produced by FT-LLaMA2 are longer, with the original judgments being, by far, the longest texts in this study, with over an order of magnitude longer texts than the *massime* and the GPT-4o summaries. The different length of the texts influences the percentage of skipped answers (as stated above) and the reading and response times, as presented in Figure 9, where every time greater than 30 minutes was removed. Notice how both GPT-4o summaries and *massime* have a similar distribution of reading times, with participants having legal knowledge spending more time on the *massime*. The original judgments and the summaries produced by FT-LLaMA2 also have a similar distribution of reading times despite a marked difference in the length of the texts.

VI. DISCUSSION

From the presented results, we can derive some information regarding the effectiveness of LLM-generated summaries in helping the understanding of legal texts. First of all, the percentage of correct answers shows that the “best” summaries, especially for the legal experts, are the *massime*. This is far

from unexpected: the *massime* are crafted by legal experts for legal experts, highlighting the important legal aspects of a judgment. When they are not present, LLM-generated summaries provide similar (FT-LLaMA2) or better (GPT-4o) percentages of correct answers. However, this analysis is not sufficient, since there is also an effect on the *incorrect* answers. There, LLM-based summaries, and GPT-4o in particular, show a strong limitation: the percentage of incorrect answers is higher. While the causes for these results cannot be derived with certainty from the questionnaires, an incorrect summary, misrepresenting the content of the original judgment, can be a reasonable culprit.

It is also important to point out how the length of the text is also an important factor in determining the answers (or their absence derived by skipping the questions). The participants required more time and had a higher probability of skipping questions for FT-LLaMA2 summaries and the original judgments. If the objective is to inform the user and help their understanding this factor should also be taken into account.

Summarising the previous observations, we can conclude that:

- *Massime* are the best possible summaries, since they are “on point”, short, and accurate. However, their generation is time-consuming and requires the availability of legal experts.
- When *massime* are not available, the summaries generated by GPT-4o are the second-best choice: they are short and provide a high percentage of understanding (as shown by the percentage of correct answers). However, they can misrepresent the legal content of the original judgment, increasing the risk of conveying incorrect information.
- FT-LLaMA2 summaries and original judgments are both of a length above a threshold after which the user is likely to ignore the content of the text, as shown by the percentage of skipped answers and the higher reading times.

Hence, GPT-4o summaries provide a good alternative when *massime* is not available, if the risk of an incorrect summary




Fig. 6. Distribution of the number of I don't know answers according to the type of text for all participants (left), participants with preexisting legal knowledge (center), and participants without legal knowledge (right).




Fig. 7. Percentage of skipped answers for each type of text for all participants (all), for participants with legal knowledge (with), and for people without legal knowledge (w/o)




Fig. 9. Boxplots showing the reading and response time for each type of text and the different categories of participants. The boxes show the 25-th and 75-th percentiles, while the whiskers denote the minimum and maximum.




Fig. 8. Boxplots of the length (in number of words) of the texts for each category. The boxes show the 25-th and 75-th percentiles, while the whiskers denote the minimum and maximum. Notice that the vertical axis is in logarithmic scale.

is taken into account. If a local LLM is used, FT-LLaMA2 summaries are, in general, not worse than the original judgment while still providing shorter summaries, even if not short enough to be considered as fast to read and understand as *massime* or GPT-4o summaries.

VII. CONCLUSION

This study investigated the capability of LLMs to generate comprehensible summaries of legal judgments issued by the Italian Constitutional Court. The main aim was to assess the efficacy of these models in rendering complex legal texts accessible to all individuals, even without preexisting legal knowledge. The results suggest that small open-source LLMs have not yet achieved the quality and comprehensibility of *massime* or GPT-4o generated summaries for Italian legal texts. While the quality of the latter is still not at the level of the former, generating summaries with GPT-4o is much cheaper and faster than producing the *massime*, representing a significant advantage in favor of LLMs. Further develop-

ment and improvements are necessary to bridge this gap and make open-source models a more effective solution for the accessibility and transparency of laws. While these models can support legal work the limitations of automatic summarization have important ethical implications that must be considered before widespread use.

For future research, solving the problem of the short context for open-source LLM is an important step. Also, integrating existing legal knowledge in the generation of the summaries could be a significant step for accuracy. Finally, an important next step would be the assistance of legal experts in evaluating the generated summaries, to obtain a more comprehensive understating of the types of issues that appear when using LLMs for summarization tasks.

REFERENCES

- [1] Flavia Acheda, David Preti, Davide Venditti, Leonardo Ranaldi, Cristina Giannone, Fabio Massimo Zanzotto, Andrea Favalli, Raniero Romagnoli, et al. Legal summarization: to each court its own model. 2023.
- [2] Benjamin Alarie, Anthony Niblett, and Albert H Yoon. How artificial intelligence will affect the practice of law. *University of Toronto Law Journal*, 68(supplement 1):106–124, 2018.
- [3] Miguel B Almeida and Andre FT Martins. Fast and robust compressive summarization with dual decomposition and multi-task learning. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 196–206, 2013.
- [4] Amos Azaria and Tom Mitchell. The internal state of an llm knows when it’s lying. *arXiv preprint arXiv:2304.13734*, 2023.
- [5] Satanjeev Banerjee and Alon Lavie. Meteor: An automatic metric for mt evaluation with improved correlation with human judgments. In *Proceedings of the acl workshop on intrinsic and extrinsic evaluation measures for machine translation and/or summarization*, pages 65–72, 2005.
- [6] Pierpaolo Basile, Elio Musacchio, Marco Polignano, Lucia Siciliani, Giuseppe Fiameni, and Giovanni Semeraro. Llamantino: Llama 2 models for effective text generation in italian language. *arXiv preprint arXiv:2312.09993*, 2023.
- [7] Bruce G Buchanan and Thomas E Headrick. Some speculation about artificial intelligence and legal reasoning. *Stan. L. Rev.*, 23:40, 1970.
- [8] Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, et al. A survey on evaluation of large language models. *ACM Transactions on Intelligent Systems and Technology*, 15(3):1–45, 2024.
- [9] Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023.
- [10] Noel Cressie and Timothy RC Read. Multinomial goodness-of-fit tests. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 46(3):440–464, 1984.
- [11] Aniket Dero, Kripabandhu Ghosh, and Saptarshi Ghosh. How ready are pre-trained abstractive models and llms for legal case judgement summarization? *arXiv preprint arXiv:2306.01248*, 2023.
- [12] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [13] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian journal of statistics*, pages 65–70, 1979.
- [14] Rahul C Kore, Prachi Ray, Priyanka Lade, and Amit Nerurkar. Legal document summarization using nlp and ml techniques. *Int. J. Eng. Comput. Sci.*, 9:25039–25046, 2020.
- [15] Chin-Yew Lin. Improving summarization performance by sentence compression—a pilot study. In *Proceedings of the sixth international workshop on Information retrieval with Asian languages*, pages 1–8, 2003.
- [16] Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pages 74–81, 2004.
- [17] Hui Lin and Jeff Bilmes. A class of submodular functions for document summarization. In *Proceedings of the 49th annual meeting of the association for computational linguistics: human language technologies*, pages 510–520, 2011.
- [18] Yang Liu. Fine-tune bert for extractive summarization. *arXiv preprint arXiv:1903.10318*, 2019.
- [19] Yang Liu and Mirella Lapata. Text summarization with pretrained encoders. *arXiv preprint arXiv:1908.08345*, 2019.
- [20] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2017.
- [21] Lauren Martin, Nick Whitehouse, Stephanie Yiu, Lizzie Catterson, and Rivindu Perera. Better call gpt, comparing large language models against lawyers. *arXiv preprint arXiv:2401.16212*, 2024.
- [22] Hajime Morita, Ryohei Sasano, Hiroya Takamura, and Manabu Okumura. Subtree extractive summarization via submodular maximization. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1023–1032, 2013.
- [23] Jianmo Ni, Gustavo Hernandez Abrego, Noah Constant, Ji Ma, Keith B Hall, Daniel Cer, and Yinfei Yang. Sentence-t5: Scalable sentence encoders from pre-trained text-to-text models. *arXiv preprint arXiv:2108.08877*, 2021.
- [24] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pages 311–318, 2002.
- [25] Guilherme Penedo, Quentin Malartic, Daniel Hesslow, Ruxandra Cojocaru, Alessandro Cappelli, Hamza Alobeidli, Baptiste Pannier, Ebtesam Almazrouei, and Julien Launay. The refinedweb dataset for falcon llm: outperforming curated corpora with web data, and web data only. *arXiv preprint arXiv:2306.01116*, 2023.
- [26] Alexander M Rush, Sumit Chopra, and Jason Weston. A neural attention model for abstractive sentence summarization. *arXiv preprint arXiv:1509.00685*, 2015.
- [27] Abigail See, Peter J Liu, and Christopher D Manning. Get to the point: Summarization with pointer-generator networks. *arXiv preprint arXiv:1704.04368*, 2017.
- [28] Gemma Team, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.
- [29] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [30] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shriti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- [31] Dingding Wang, Tao Li, Shenghuo Zhu, and Chris Ding. Multi-document summarization via sentence-level semantic analysis and symmetric matrix factorization. In *Proceedings of the 31st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 307–314, 2008.
- [32] Qiang Wang, Bei Li, Tong Xiao, Jingbo Zhu, Changliang Li, Derek F Wong, and Lidia S Chao. Learning deep transformer models for machine translation. *arXiv preprint arXiv:1906.01787*, 2019.
- [33] David M Zajic, Bonnie Dorr, Jimmy Lin, and Richard Schwartz. Sentence compression as a component of a multi-document summarization system. In *Proceedings of the 2006 document understanding workshop, New York*, 2006.
- [34] Jingqing Zhang, Yao Zhao, Mohammad Saleh, and Peter Liu. Pegasus: Pre-training with extracted gap-sentences for abstractive summarization. In *International conference on machine learning*, pages 11328–11339. PMLR, 2020.
- [35] Ming Zhong, Pengfei Liu, Yiran Chen, Danqing Wang, Xipeng Qiu, and Xuanjing Huang. Extractive summarization as text matching. *arXiv preprint arXiv:2004.08795*, 2020.
- [36] Zolt Zódi. The limits of plain legal language: understanding the comprehensible style in law. *International Journal of Law in Context*, 15(3):246–262, 2019.